

# LAB\_8

I used the **terminal** to work with the **postgresql**, so here will be the screens of the terminal

## Exercise 1

- Now we do not have additional indecies except for the one for the primary key

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c '\d customer"
Expanded display is on.
You are now connected to database "customers" as user "postgres".
      Table "public.customer"
  Column | Type   | Collation | Nullable | Default
-----+-----+-----+-----+-----
 id      | integer |           | not null |
 name    | text    |           | not null |
 address | text    |           | not null |
 review  | text    |           |          |
Indexes:
    "customer_pkey" PRIMARY KEY, btree (id)
```

### 1. 1st query analysis:

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "EXPLAIN (ANALYSE) SELECT name FROM customer GROUP BY name ORDER BY name;"
Expanded display is on.
You are now connected to database "customers" as user "postgres".
-[ RECORD 1 ]-----
QUERY PLAN | Sort (cost=29485.04..29680.94 rows=78360 width=14) (actual time=1642.667..1894.425 rows=120120 loops=1)
-[ RECORD 2 ]-----
QUERY PLAN |      Sort Key: name
-[ RECORD 3 ]-----
QUERY PLAN |      Sort Method: external merge  Disk: 2864kB
-[ RECORD 4 ]-----
QUERY PLAN |      -> HashAggregate (cost=20378.50..23115.22 rows=78360 width=14) (actual time=361.634..504.781 rows=120120 loops=1)
-[ RECORD 5 ]-----
QUERY PLAN |      Group Key: name
-[ RECORD 6 ]-----
QUERY PLAN |      Planned Partitions: 4  Batches: 5  Memory Usage: 4145kB  Disk Usage: 3688kB
-[ RECORD 7 ]-----
QUERY PLAN |      -> Seq Scan on customer (cost=0.00..8066.00 rows=200000 width=14) (actual time=0.026..66.446 rows=200000 loops=1)
-[ RECORD 8 ]-----
QUERY PLAN | Planning Time: 1.352 ms
-[ RECORD 9 ]-----
QUERY PLAN | Execution Time: 1913.452 ms
```

$Cost_{1\_no\_index} = 60,862.16$

$Execution\_time_{1\_no\_index} = 1913.452 (ms)$

### 2. 2nd query analysis

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "EXPLAIN (ANALYSE) SELECT name, address FROM customer GROUP BY name, address ORDER BY name, address;"
Expanded display is on.
You are now connected to database "customers" as user "postgres".
-[ RECORD 1 ]-----
QUERY PLAN | Group (cost=33197.14..34697.14 rows=200000 width=58) (actual time=2006.128..2907.515 rows=200000 loops=1)
-[ RECORD 2 ]-----
QUERY PLAN |      Group Key: name, address
-[ RECORD 3 ]-----
QUERY PLAN |      -> Sort (cost=33197.14..33697.14 rows=200000 width=58) (actual time=2006.118..2753.343 rows=200000 loops=1)
-[ RECORD 4 ]-----
QUERY PLAN |      Sort Key: name, address
-[ RECORD 5 ]-----
QUERY PLAN |      Sort Method: external merge  Disk: 13528kB
-[ RECORD 6 ]-----
QUERY PLAN |      -> Seq Scan on customer (cost=0.00..8066.00 rows=200000 width=58) (actual time=0.021..123.326 rows=200000 loops=1)
-[ RECORD 7 ]-----
QUERY PLAN | Planning Time: 1.571 ms
-[ RECORD 8 ]-----
QUERY PLAN | Execution Time: 2937.506 ms
```

$Cost_{2\_no\_index} = 70,460.28$

$Execution\_time_{2\_no\_index} = 2937.506 (ms)$

### 3. 3rd query analysis

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "EXPLAIN (ANALYSE) SELECT name, address, review FROM customer GROUP BY name, address, review ORDER BY name;"
Expanded display is on.
You are now connected to database "customers" as user "postgres".
-[ RECORD 1 ]-----
QUERY PLAN | Group (cost=22973.91..44294.63 rows=200000 width=207) (actual time=731.392..1809.973 rows=200000 loops=1)
-[ RECORD 2 ]-----
QUERY PLAN | Group Key: name, address, review
-[ RECORD 3 ]-----
QUERY PLAN | -> Gather Merge (cost=22973.91..43044.64 rows=166666 width=207) (actual time=731.387..1599.916 rows=200000 loops=1)
-[ RECORD 4 ]-----
QUERY PLAN | Workers Planned: 2
-[ RECORD 5 ]-----
QUERY PLAN | Workers Launched: 2
-[ RECORD 6 ]-----
QUERY PLAN | -> Group (cost=21973.89..22807.22 rows=83333 width=207) (actual time=714.630..1052.051 rows=66667 loops=3)
-[ RECORD 7 ]-----
QUERY PLAN | Group Key: name, address, review
-[ RECORD 8 ]-----
QUERY PLAN | -> Sort (cost=21973.89..22182.22 rows=83333 width=207) (actual time=714.617..978.321 rows=66667 loops=3)
-[ RECORD 9 ]-----
QUERY PLAN | Sort Key: name, address, review
-[ RECORD 10 ]-----
QUERY PLAN | Sort Method: external merge Disk: 14656kB
-[ RECORD 11 ]-----
QUERY PLAN | Worker 0: Sort Method: external merge Disk: 14088kB
-[ RECORD 12 ]-----
QUERY PLAN | Worker 1: Sort Method: external merge Disk: 14360kB
-[ RECORD 13 ]-----
QUERY PLAN | -> Parallel Seq Scan on customer (cost=0.00..6899.33 rows=83333 width=207) (actual time=0.023..51.960 rows=66667 loops=3)
-[ RECORD 14 ]-----
QUERY PLAN | Planning Time: 1.643 ms
-[ RECORD 15 ]-----
QUERY PLAN | Execution Time: 1843.406 ms
```

*Cost\_3\_no\_index* = 94,933.41

*Execution\_time\_3\_no\_index* = 1843.406 (ms)

- I created the index na\_key

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "CREATE INDEX IF NOT EXISTS n_key ON customer USING btree (name);"
could not change directory to "/home/andrew": Permission denied
Expanded display is on.
You are now connected to database "customers" as user "postgres".
CREATE INDEX
```

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "\d customer"
could not change directory to "/home/andrew": Permission denied
Expanded display is on.
You are now connected to database "customers" as user "postgres".
Table "public.customer"
Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----
id      | integer |          | not null |
name    | text    |          | not null |
address | text    |          | not null |
review  | text    |          |          |
Indexes:
    "customer_pkey" PRIMARY KEY, btree (id)
    "n_key" btree (name)
```

### 1. 1st query analysis

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "EXPLAIN (ANALYSE) SELECT name FROM customer GROUP BY name ORDER BY name;"
could not change directory to "/home/andrew": Permission denied
Expanded display is on.
You are now connected to database "customers" as user "postgres".
-[ RECORD 1 ]-----
QUERY PLAN | Group (cost=0.42..10356.33 rows=78360 width=14) (actual time=0.164..242.383 rows=120120 loops=1)
-[ RECORD 2 ]-----
QUERY PLAN | Group Key: name
-[ RECORD 3 ]-----
QUERY PLAN | -> Index Only Scan using n_key on customer (cost=0.42..9856.33 rows=200000 width=14) (actual time=0.160..143.735 rows=200000 loops=1)
-[ RECORD 4 ]-----
QUERY PLAN | Heap Fetches: 36763
-[ RECORD 5 ]-----
QUERY PLAN | Planning Time: 1.640 ms
-[ RECORD 6 ]-----
QUERY PLAN | Execution Time: 255.695 ms
```

*Cost\_1\_with\_index* = 20,212.66

*Execution\_time\_1\_with\_index* = 255.695 (ms)

## 2. 2nd query analysis

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "EXPLAIN (ANALYSE) SELECT name, address FROM customer GROUP BY name, address ORDER BY name, address;"
[sudo] password for andrew:
could not change directory to "/home/andrew": Permission denied
Expanded display is on.
You are now connected to database "customers" as user "postgres".
-[ RECORD 1 ]-----
QUERY PLAN | Group (cost=33197.14..34697.14 rows=200000 width=58) (actual time=2064.388..2974.152 rows=200000 loops=1)
-[ RECORD 2 ]-----
QUERY PLAN | Group Key: name, address
-[ RECORD 3 ]-----
QUERY PLAN | -> Sort (cost=33197.14..33697.14 rows=200000 width=58) (actual time=2064.377..2816.890 rows=200000 loops=1)
-[ RECORD 4 ]-----
QUERY PLAN | Sort Key: name, address
-[ RECORD 5 ]-----
QUERY PLAN | Sort Method: external merge Disk: 13528kB
-[ RECORD 6 ]-----
QUERY PLAN | -> Seq Scan on customer (cost=0.00..8066.00 rows=200000 width=58) (actual time=0.469..169.813 rows=200000 loops=1)
-[ RECORD 7 ]-----
QUERY PLAN | Planning Time: 2.520 ms
-[ RECORD 8 ]-----
QUERY PLAN | Execution Time: 3003.553 ms
```

$Cost\_2\_with\_index = 70,460.28$

$Execution\_time\_2\_with\_index = 3003.553\ (ms)$

## 3. 3rd query analysis

```
> sudo -u postgres psql -c '\x' -c '\c customers' -c "EXPLAIN (ANALYSE) SELECT name, address, review FROM customer GROUP BY name, address, review ORDER BY name;"
could not change directory to "/home/andrew": Permission denied
Expanded display is on.
You are now connected to database "customers" as user "postgres".
-[ RECORD 1 ]-----
QUERY PLAN | Group (cost=0.84..38382.27 rows=200000 width=207) (actual time=1.256..1583.542 rows=200000 loops=1)
-[ RECORD 2 ]-----
QUERY PLAN | Group Key: name, address, review
-[ RECORD 3 ]-----
QUERY PLAN | -> Incremental Sort (cost=0.84..36882.27 rows=200000 width=207) (actual time=1.251..1426.849 rows=200000 loops=1)
-[ RECORD 4 ]-----
QUERY PLAN | Sort Key: name, address, review
-[ RECORD 5 ]-----
QUERY PLAN | Presorted Key: name
-[ RECORD 6 ]-----
QUERY PLAN | Full-sort Groups: 5991 Sort Method: quicksort Average Memory: 36kB Peak Memory: 36kB
-[ RECORD 7 ]-----
QUERY PLAN | Pre-sorted Groups: 149 Sort Method: quicksort Average Memory: 26kB Peak Memory: 32kB
-[ RECORD 8 ]-----
QUERY PLAN | -> Index Scan using n_key on customer (cost=0.42..29659.91 rows=200000 width=207) (actual time=0.083..602.326 rows=200000 loops=1)
-[ RECORD 9 ]-----
QUERY PLAN | Planning Time: 3.030 ms
-[ RECORD 10 ]-----
QUERY PLAN | Execution Time: 1605.376 ms
```

$Cost\_3\_no\_index = 104,924.45$

$Execution\_time\_3\_with\_index = 1605.376\ (ms)$

## Analysing of the result

### • 1st query

$Cost\_1\_no\_index = 60,862.16 > Cost\_1\_with\_index = 20,212.66$

**Since we created the index for the column we are work with, we have the smaller result after the creating the index**

$Execution\_time\_1\_no\_index = 1913.452\ (ms) > Execution\_time\_1\_with\_index = 255.695\ (ms)$

### • 2nd query

$Cost\_2\_no\_index = 70,460.28 = Cost\_2\_with\_index = 70,460.28$

**Here we see that the result for the cost is the same, so we could even not use the index**

$Execution\_time\_2\_no\_index = 2937.506\ (ms) < Execution\_time\_2\_with\_index = 3003.553\ (ms)$

### • 3rd query

$Cost\_3\_no\_index = 94,933.41 < Cost\_3\_with\_index = 104,924.45$

**The result says that in this case it would be better not to use the index for this query**

$Execution\_time\_3\_no\_index = 1843.406\ (ms) > Execution\_time\_3\_with\_index = 1605.376\ (ms)$

## Exercise 2

### result

### • 1st query

```
SELECT film.film_id, film.title from film, category, film_category where film.film_id NOT IN (SELECT
inventory.film_id From inventory) and film.film_id=film_category.film_id and
category.category_id=film_category.category
y_id and (category.name='Horror' or category.name='Sci-fi') and (film.rating='R' or film.rating='PG-13')
```

```

> sudo -u postgres psql -c '\c dvdrental' -c "EXPLAIN ANALYSE SELECT film.film_id, film.title from film, category, film_category wh
entory.film_id From inventory) and film.film_id=film_category.film_id and category.category_id=film_category.category_id and (catego
Sci-fi') and (film.rating='R' or film.rating='PG-13'))"
You are now connected to database "dvdrental" as user "postgres".

                                QUERY PLAN
-----
Hash Join  (cost=104.40..176.84 rows=23 width=19) (actual time=5.523..6.809 rows=2 loops=1)
  Hash Cond: (film.film_id = film_category.film_id)
    -> Seq Scan on film  (cost=82.26..153.76 rows=187 width=19) (actual time=4.541..5.888 rows=16 loops=1)
          Filter: ((NOT (hashed SubPlan 1)) AND ((rating = 'R'::mpaa_rating) OR (rating = 'PG-13'::mpaa_rating)))
          Rows Removed by Filter: 984
          SubPlan 1
            -> Seq Scan on inventory  (cost=0.00..70.81 rows=4581 width=2) (actual time=0.014..1.754 rows=4581 loops=1)
    -> Hash  (cost=20.58..20.58 rows=125 width=2) (actual time=0.739..0.744 rows=56 loops=1)
          Buckets: 1024  Batches: 1  Memory Usage: 10kB
          -> Hash Join  (cost=1.26..20.58 rows=125 width=2) (actual time=0.106..0.700 rows=56 loops=1)
                Hash Cond: (film_category.category_id = category.category_id)
                  -> Seq Scan on film_category  (cost=0.00..16.00 rows=1000 width=4) (actual time=0.020..0.275 rows=1000 loops=1)
                  -> Seq Scan on category  (cost=1.24..1.24 rows=2 width=4) (actual time=0.045..0.047 rows=1 loops=1)
                        Buckets: 1024  Batches: 1  Memory Usage: 9kB
                        -> Seq Scan on category  (cost=0.00..1.24 rows=2 width=4) (actual time=0.029..0.034 rows=1 loops=1)
                              Filter: (((name)::text = 'Horror'::text) OR ((name)::text = 'Sci-fi'::text))
                              Rows Removed by Filter: 15

Planning Time: 4.380 ms
Execution Time: 7.245 ms
(19 rows)

```

The most expensive operation was: **Scannin on film**

- **2nd query**

```
SELECT DISTINCT ON (V2.city_id) * FROM (SELECT V.city_id, V.staff_id, SUM(amount) as total FROM (payment
INNER JOIN customer USING(customer_id) INNER JOIN address USING(address_id)) V WHERE date_part('month',
age(date('2007-05-14 13:44:29.996577'),date(V.payment_date))) <= 1 GROUP BY V.city_id, V.staff_id) V2 ORDER BY
V2.city_id, V2.total ASC
```

```

psql -u postgres psql -c "\c dvdrental" -c EXPLAIN ANALYSE SELECT DISTINCT ON (V2.city_id) * FROM (SELECT V.city_id, V.staff_id, SUM(amount) as total FROM (payment INNER JOIN customer USING(customer_id, age(date('2007-05-14 13:44:29.996577')), date(V.payment_date))) <= 1 GROUP BY V.city_id, V.staff_id) V2 ORDER BY V2.city_id, V2.total ASC
You are now connected to database "dvdrental" as user "postgres".

QUERY PLAN
-----
Unique (cost=667.36..673.35 rows=200 width=36) (actual time=82.563..83.176 rows=597 loops=1)
  -> Sort (cost=667.36..670.36 rows=1198 width=36) (actual time=82.561..82.762 rows=1194 loops=1)
    Sort Key: address.city_id, (sum(payment.amount))
    Sort Method: quicksort Memory: 104kB
    -> HashAggregate (cost=579.15..594.13 rows=1198 width=36) (actual time=79.486..80.987 rows=1194 loops=1)
      Group Key: address.city_id, payment.staff_id
      Batches: 1 Memory Usage: 577kB
      -> Hash Join (cost=44.05..542.66 rows=4865 width=10) (actual time=6.867..62.913 rows=11309 loops=1)
        Hash Cond: (customer.address_id = address.address_id)
        -> Hash Join (cost=22.48..508.24 rows=4865 width=10) (actual time=6.088..54.036 rows=11309 loops=1)
          Hash Cond: (payment.customer_id = customer.customer_id)
          -> Seq Scan on payment (cost=0.00..472.90 rows=4865 width=10) (actual time=4.923..40.598 rows=11309 loops=1)
            Filter: (date_part('month':text, age('2007-05-14'))::date)::timestamp with time zone, (date(payment_date))::timestamp with time zone)) <= '1':double precision)
            Rows Removed by Filter: 3287
          -> Hash (cost=14.99..14.99 rows=599 width=6) (actual time=1.065..1.067 rows=599 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 31kB
            -> Seq Scan on customer (cost=0.00..14.99 rows=599 width=6) (actual time=0.022..0.530 rows=599 loops=1)
          -> Hash (cost=14.03..14.03 rows=603 width=6) (actual time=0.714..0.715 rows=603 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 32kB
            -> Seq Scan on address (cost=0.00..14.03 rows=603 width=6) (actual time=0.022..0.368 rows=603 loops=1)
Planning Time: 5.449 ms
Execution Time: 83.811 ms
(22 rows)

```

The most expensive operation was: **sort**